

Algorithmic Analysis

1 < log(log n) < log n < log^2 n ≤ √n < n < n log n < n^2 < 2^n < 3^n < n! < n^n

T(n) = T(n-1) + O(1)	O(n)
T(n) = T(n-1) + O(logn)	O(n log n)
T(n) = T(n-1) + O(n)	O(n^2)
T(n) = 2T(n-1) + O(1)	O(2^n)
T(n) = T(n/2) + O(1)	O(log n)
T(n) = T(n/2) + O(n)	O(n)
T(n) = 2T(n/2) + O(1)	O(n log n)
T(n) = 2T(n/2) + O(n)	O(n)
T(n) = 2T(n/2) + O(n^2)	O(n^2)
T(n) = T(sqrt(n)) + O(1)	O(log(log n))

Searching:

Binary Search

```
int search(A, key, n)
begin = 0
end = n-1
while begin < end do:
    mid = begin + (end-begin)/2;
    if key <= A[mid] then
        end = mid
    else begin = mid+1
return (A[begin]==key) ? begin : -1
Loop invariants:
Correctness - A[begin] <= key <= A[end]
Performance - (end-begin) <= n/2k in iteration k.
```

Peak Finding

```
FindPeak(A, n)
if A[n/2+1] > A[n/2]:
    FindPeak (A[n/2+1..n], n/2)
else if A[n/2-1] > A[n/2]:
    FindPeak (A[1..n/2-1], n/2)
else A[n/2] is a peak; return n/2
```

Sorting:

```
BubbleSort(A, n)
repeat (n times/until no swaps):
    for j in [1...n-1]:
        if A[j] > A[j+1] then swap(A[j], A[j+1])
```

```
SelectionSort(A, n)
for j in [1...n-1]:
    find minimum element A[k] in A[j..n]
    swap(A[j], A[k])
```

```
InsertionSort(A, n)
for j in [2...n]:
    key = A[j]
    i = j-1
    while (i > 0) and (A[i] > key):
        A[i+1] = A[i]
        i = i-1
    A[i+1] = key
```

```
MergeSort(A, n)
if (n=1) then return;
else:
    X ~MergeSort(A[1..n/2], n/2);
    Y ~MergeSort(A[n/2+1, n], n/2);
return Merge (X,Y, n/2);
```

master theorem

$$T(n) = aT(\frac{n}{b}) + f(n) \quad a \geq 0, b > 1$$
$$= \begin{cases} \Theta(n^{\log_b a}) & \text{if } f(n) < n^{\log_b a} \text{ polynomially} \\ \Theta(n^{\log_b a} \log n) & \text{if } f(n) = n^{\log_b a} \\ \Theta(f(n)) & \text{if } f(n) > n^{\log_b a} \text{ polynomially} \end{cases}$$

```
QuickSort(A[1..n], n)
if (n==1) then return;
else:
    Choose plndex (random until >1/10 for paranoid)
    p = partition(A[1..n], n, plndex)
    x = QuickSort(A[1..p-1], p-1)
    y = QuickSort(A[p+1..n], n-p)
partition(A[1..n], n, plndex):
    // Assume no duplicates, n>1
    pivot = A[plndex]; // plndex is the index of pivot
    swap(A[1], A[plndex]); // store pivot in A[1]
    low = 2; // start after pivot in A[1]
    high = n+1; // Define: A[n+1] = infinity
    while (low < high)
        while (A[low] < pivot) and (low < high) do low++;
        while (A[high] > pivot) and (low < high) do high--;
        if (low < high) then swap(A[low], A[high]);
    swap(A[1], A[low-1]);
    return low-1;
For duplicates:
3-way partitioning (2 pivots)
1-pass, maintain 4 array regions
```

< pivot	= pivot	In Progress	> pivot
---------	---------	-------------	---------

>2 pivots still same time complexity

Order Statistics


```
Finds the kth smallest element in an array. O(n)
QuickSelect(A[1..n], n, k)
if (n == 1) then return A[1];
else:
    Choose random pivot index plndex (paranoid)
    p = partition(A[1..n], n, plndex)
    if (k == p) then return A[p];
    else if (k < p):
        return QuickSelect(A[1..p-1], k)
    else if (k > p):
        return QuickSelect(A[p+1], k - p)
```

Node type	#Keys		#Children	
	Min	Max	Min	Max
Root	1	b-1	2	b
Internal	a-1	b-1	a	b
Leaf	a-1	b-1	0	0

Name	Best	Avg	Worst	+Space	Stable	Invariant
BubbleSort	O(n)	O(n^2)	O(n^2)	O(1)	Yes	Largest j items sorted in correct positions
SelectionSort	O(n^2)	O(n^2)	O(n^2)	O(1)	No	Smallest j items sorted in correct positions
InsertionSort	O(n)	O(n^2)	O(n^2)	O(1)	Yes	First j items sorted, may not be correct
MergeSort	O(nlogn)	O(nlogn)	O(nlogn)	O(n)	Yes	Sort within buckets before merging
QuickSort	O(nlogn)	O(nlogn)	O(n^2) (non-random pivot)	O(1)	No	Pivot in correct position Left partition all smaller than pivot Right partition all larger than pivot
HeapSort	O(nlogn)	O(nlogn)	O(nlogn)	O(n)	No	-

Trees

- For leaf node v: h(v) = 0, For root/internal node v: h(v) = max(h(v.left), h(v.right)) + 1
- Keep items in a hierarchical order to facilitate efficient search
- Keep items sorted in inorder traversal sequence
- Binary Search Tree (BST)
- Not necessarily balanced unless h = O(log n)
- log n - 1 ≤ h ≤ n
- On a balanced BST: all operations run in O(log n) time
- AVL Tree
- Each node stores absolute height or relative height to parent
- On insert & delete update height
- A node v is height-balanced if: |v.left.height - v.right.height| ≤ 1
- 2^h/2 ≤ n < 2^h
- Insert: O(logn)
 - Find location O(logn)
 - Walk up tree O(logn)
 - checkBalance O(1) -> unbalanced -> max 2 rotations -> O(1)



- Delete: O(logn)
 - Find location O(logn)
 - If v has 2 children, swap with successor
 - Delete v and reconnect children
 - For every ancestor of v: check height balance & rotate -> O(logn)

Trie

- Each node stores character & boolean endFlag
- Space: O(text size)
- Search, insert, delete: O(L)
- (a,b)-Tree
- Each internal node: a ≤ children.length < b, where 2 ≤ a ≤ (b+1)/2
- Each root/internal node has children.length = keys.length + 1
- Child has key range (v_i-1, v_i]
- All leaf nodes have the same depth from root
- Maximum height: O(log_a n) + 1
- Minimum height: O(log_b n)
- B-trees are simply (B, 2B)-trees
- Split: promote median of overfilled keys to parent key (proactive/passive)
- Delete: if leaf just delete, else swap with pred/succ then delete (proactive/passive)
- Pred: rightmost key in left subtree, Succ: leftmost key in right subtree
- Merge: Demote prev key of parent node to left sibling (propagate up)
- Share: Merge then Split (if overfilled)

Augmented Trees CS2040S 21/22 Notes by Justin Peng Methodology:

- Choose underlying data structure (tree, hash table, linked list, stack, etc.)
- Determine additional info needed.
- Modify data structure to maintain additional info when the structure changes. (subject to insert/delete/etc.)
- Develop new operations

Dynamic Order Statistics

Support insert, delete, select(k): find the kth smallest node

- AVL Tree
- Weight of each node
- Update weights
- Select, Rank

Interval Queries

Support insert, delete, query(x): find an interval that contains x

- AVL Tree
- Upper bound of interval, upper bound of subtree
- Update upper bound of subtree
- Query: O(logn)

Orthogonal Range Searching

Support insert, delete, query(1Dinterval): find all points in 1Dinterval

- BST
- Store all points in leaves of the tree, internal node stores max of left subtree
- Update max of left subtree
- Query, FindSplit, LeftTraversal, RightTraversal, AllTraversal
 - LeftTraversal: either AllTraversal right & recurse left, or recurse right -> O(logn)
 - RightTraversal: either AllTraversal left & recurse right, or recurse left -> O(logn)
 - AllTraversal: O(k), k = no. of items found
 - Query: O(k + logn)
 - BuildTree: O(nlogn)
 - Space complexity: O(n)

If just need to count no. of points, store and use weight of each node instead of AllTraversal

Support insert, delete, query(2Dinterval): find all points in 2Dinterval

- 1D Range Tree for x-coords
- 1D Range Tree for y-coords in each node
- Same as above
- Same as above
 - FindSplit: O(logn)
 - Traversal on x: O(logn)
 - Traversal on y: O(logn) (for each x traversal)
 - Output: O(k)
 - Query: O(log^2 n + k)
 - BuildTree: O(nlogn)
 - Space complexity: O(nlogn)

kd-Trees

- Alternate levels in the tree
- Each level divides the points in the plane in half
- Supports more efficient updates
- Often better in practice
- Good for a variety of queries (e.g., nearest neighbor)

Type	Search	Insert	Delete	Pred/Succ	Traversal
BST	O(h)	O(h)	O(h)	O(h)	O(h)
AVL	O(log n)	O(log n)	O(log n)	O(log n)	O(n)
(a,b)	O(log_a n)	O(blog_a n)	O(log n)		
Trie	O(L)	O(L)	O(L)		O(n)
HashTable	O(1) or O(n)	O(1)	O(1)	--	--

- Hashing**
 - h = hash function
 - n = no. of keys
 - m = no. of buckets
- Hash Functions**
 - h: $U \rightarrow \{1...m\}$ $O(1)$
 - Store key k in bucket h(k)
 - Collisions are inevitable (pigeonhole principle)
- Chaining (Collision Resolution)**
 - Each bucket \rightarrow linked list
 - Space: $O(m + n)$
 - Insert: $O(1)$
 - Calculate h(key) $\rightarrow O(1)$
 - Get bucket h(key) $\rightarrow O(1)$
 - Add (key, value) to linked list $\rightarrow O(1)$
- Search: $O(n)$ worst, $O(1)$ best, $O(1 + n/m)$ SUHA
 - Calculate h(key) $\rightarrow O(1)$
 - Get bucket h(key) $\rightarrow O(1)$
 - Get items in linked list $\rightarrow O(n)$
- Delete: $O(n)$ worst, $O(1)$ best, $O(1 + n/m)$ SUHA
- Inserting n elements: max cost $O(\log n / \log(\log n))$
- Simple Uniform Hashing Assumption (SUHA)**
 - Every key equally likely to map to every bucket
 - Keys mapped independently
 - load(HashTable) = n/m = avg # items per bucket
 - Search: $O(1 + n/m) = O(1)$ if load < 1
- Open Addressing (Collision Resolution)**
 - On collision, probe sequence of buckets (all) till empty found
- h(key, i) : $U \rightarrow \{1..m\}$ where i = no. of collisions (enumerate all buckets)
- Linear Probing: check next bucket
 - Problem: clusters of size $O(\log n)$ if table is 1/4 full (not $O(1)$)
 - Reality: caching \rightarrow fast access nearby elems \rightarrow faster
- Search: check each bucket h(key, i) from 1 to m
- Delete: set removed bucket as tombstone value for search
- n items, m buckets, SUHA \rightarrow expected cost $\leq 1/(1-\alpha)$
- Sensitive to hash function choice & load; performance degrades as $\alpha \rightarrow 1$
- Double Hashing**
 - $h(k, i) = (ff(k) + i * g(k)) \bmod m$
 - if g(k) relatively prime to m \rightarrow h(k, i) enumerate all buckets
- Table Resizing**
 - If $(n == m)$, then $m = 2m$; If $(n < m/4)$, then $m = m/2$
 - Only grow when at least $m/2$ new items added
 - Only shrink when at least $m/4$ items deleted
 - Every time you shrink a table of size m, at least $m/4$ items were deleted
 - Growing: $O(n + 2n + n) = O(n)$ assuming $m < n < 2m$ (new table size)
 - Scanning old HT: $O(m) = O(n)$, init new HT: $O(2m) = O(n)$
 - Inserting each element in new hash table: $O(1)$
- Amortized expected cost for operations: $O(1)$ (most inserts take 1)
 - Amortized cost T(n): cost of (int) k operations is $\leq k * T(n)$

- Graphs**
 - Vertices + Edges (connect 2 nodes, unique)
 - Multigraph: ≥ 1 edge between 2 nodes
 - Hypergraph: edge connects ≥ 2 nodes
 - Graph $G < V, E >; V > 0, E \subseteq \{(v, w) : (v \in V), (w \in V)\}$
- Terminology**
 - (Simple) Path: set of edges connecting 2 nodes
 - (Dis)Connected: (not)every pair of nodes connected by path
 - Cycle: form a loop
 - Tree: connected acyclic graph
 - Forest: acyclic graph
 - Degree of node: no. of adjacent/ingoing/outgoing edges
 - Degree of graph: max degree of nodes in graph
 - Diameter: shortest path dist btw furthest nodes
 - Sparse: $E = O(V)$; Dense: $E = O(V^2)$
- Special Graphs**
 - Star: 1 central node; all edges adjacent to center (deg n-1, diam 2)
 - Clique (complete graph): all node pairs connected by edges (deg n-1, diam 1)
 - Line/path: line of nodes (deg 2, diam n-1)
 - Cycle: circle of nodes (deg 2, diam n/2 or n/2-1)
 - Bipartite: 2 sets of nodes, no edges within same set (deg ?, diam ?)
- Representations**
 - Adjacency List (Sparse): size V array of neighbour LinkedLists $O(V+E)$
 - Adjacency Matrix (Dense): $V * V$ matrix of edges $O(1)$ $O(V^2)$
 - $A^n \rightarrow$ no. of n-length walks

Query	v&w nbrs?	Enumerate nbrs	Find any nbr
A-List	Slow	Fast	Fast
A-Matrix	Fast	Slow	Slow

- BFS**
 - Queue frontier, boolean[] visited, int[] parent; frontier.add(start)
 - While frontier not empty, v = frontier.poll()
 - Enum v.nbrs, if not visited then set parent & add to frontier
 - Take parents of end until start to get shortest path
- $O(V+E)$, shortest path, adjacency list, form tree
- SSSP for weightless/same weight edges (min hops not distance)
- DFS**
 - Stack s, boolean[] visited, s.push(start)
 - While stack not empty, v = frontier.pop()
 - Enum v.nbrs, if not visited then push to s
- $O(V+E)$, not shortest path, adjacency list, form tree
- Directed Graphs**
 - Edges are directed (order matters, start & end nodes)
 - In/Out-Degree: no. of incoming/outgoing edges
 - BFS/DFS work, only consider outgoing edges as nbrs
- Topological Sort**
 - Topo Order: sequential total order of all nodes, edges point fwd (DAG)
 - Sort: post-order DFS \rightarrow process node on last visit (after children)
 - Process node: prepend to T.O. (leaves processed first then parents)
 - Time complexity: same as DFS $O(V+E)$
 - Not unique
 - Kahn's algo: add all roots to T.O., remove roots, repeat. $O(V+E)$
- Connected Components**
 - Undirected: path from v to w
 - Strongly CC (digraph): path from v to w & w to v
 - Graph of Strongly CC is acyclic
- Bellman-Ford**
 - SSSP on digraph
 - relax(u, v): if (dist[v] > dist[u] + weight(u,v)) { reduce dist[v] }
 - Relax all edges: $O(E)$, V times: $O(V) \rightarrow$ total $O(VE)$
 - Early termination: entire iter of relax ops \rightarrow no change
 - Invariants: after n iters, n hop estimate is correct. $est[n] \geq dist[n]$
 - Detect negative weight cycle: V+1 iters still changing, unsolvable

- Dijkstra**
 - SSSP on digraph, no negative weights, faster than Bellman-Ford
 - Store ests in PQ, take min, add to tree & relax all outgoing
 - $O(E \log V)$ = insert/delete from PQ V times: $O(V \log V)$ + relax/decreaseKey E times: $O(E \log V)$
 - Invariant: every dequeued vertex has correct est
 - Early term: destination done once dequeued
- DAG TopoSort**
 - Relax edges in T.O., $O(E) = O(V+E)$
 - Reverse post-order DFS
 - Longest path: negate weights and find shortest path
- Union-Find**
 - Union: connect 2 nodes; Find: path connecting the 2 exists?
 - Quick-Find: Maintain array of componentIDs
 - U: update all CIDs of component $\rightarrow O(n)$
 - F: check matching CID $\rightarrow O(1)$
 - Quick-Union: Maintain array of parents
 - U: set root1 parent of root2 $\rightarrow O(n)$ (unbalanced tree)
 - F: check matching roots $\rightarrow O(n)$ (unbalanced tree)
 - Weighted Union: maintain size & parents
 - Balance tree by setting heavy root parent of light root
 - U: $O(\log n)$, F: $O(\log n)$
 - Alts: union by rank = log(size) or height
- Path Compression**
 - When root found, set parent of traversed nodes to root
 - OR set parent to grandparent (alternate nodes)
 - U: $O(\log n)$, F: $O(\log n)$
- WU + PC: m U/F ops on n objects $\rightarrow O(n + m * \alpha(m, n))$
 - Inverse Ackermann function ≤ 5
 - U: $O(\alpha(m, n))$, F: $O(\alpha(m, n))$
- MSTs**
 - Acyclic subset of edges + connects all nodes + minimum weight
 - V-1 edges; Not unique; Not the same as shortest path
 - Reweighting/negative weights are OK
 - Properties:
 - No cycles
 - Cut MST \rightarrow 2 MSTs
 - Max edge of a cycle not in MST (min edge may/not)
 - Min edge across any cut is in MST
 - Weightless/same weight graph: BFS/DFS, $O(E)$
 - Directed MST: only if acyclic + 1 root \rightarrow add min incoming edge for all nodes except root $O(E)$
 - MaxST: reverse Kruskal OR negate all edge weights + MST algo
- Generic MST Algo (Greedy)**
 - Red rule: Cycle with no red \rightarrow max edge red
 - Blue rule: Cut with no blue across \rightarrow min edge blue
 - Apply any rule to any edge until all coloured
 - Blue edges \rightarrow MST (red rule \rightarrow acyclic/tree, blue rule \rightarrow spanning)
- Prim**
 - PQ (key: min edge) for nodes across cut of {S, V-S}, S = set of nodes in MST, parent HashMap (depending on MST representation)
 - While PQ not empty, remove min node and add to S, update PQ
 - Proof: each added edge is min on some cut
 - $O(E \log V) = O(V \log V) + O(E \log V)$; like Dijkstra but weight not dist
 - Variant: edges have fixed weight range (e.g. 1-10)
 - Fixed size array of LinkedLists w/ nodes of each weight
 - $O(E) = V * O(1)$ add/remove PQ + $E * O(1)$ decreaseKey
- Kruskal**
 - Union-Find for blue tree components, sorted array of edge weights
 - Edges in ascending weight order: if join 2 nodes in same blue tree, red (max weight in cycle); else blue (min weight across cut)
 - $O(E \log V) = O(E \log E)$ sorting + $E * O(\log V)$ OR a) U/F ops
 - $\log E = O(\log V)$ since $E = O(V^2) \rightarrow \log E = O(2 \log V)$
 - Variant: edges have fixed weight range (e.g. 1-10)
 - Fixed size array of LinkedLists w/ edges of each weight
 - $O(\alpha E) = O(E)$ fill array + $E * O(\alpha)$ U/F ops

- Boruvka** CS2040S 21/22 Notes by Justin Peng
 - Create V connected components (1 per node) $O(V)$
 - Repeat: $O(\log V)$ times
 - For each CC, add min outgoing edge (BFS/DFS) $O(V+E)$
 - Merge at least k/2 components where k = no. of CC $O(V)$
 - $O((E+V) \log V) = O(E \log V)$
- Dynamic Programming**
 - Optimal sub-structure; subproblems overlap (memoise), no overlap (D&C)
 - Top-down: recurse + memoise; bottom-up: solve small + combine
 - Table view: rows = steps; columns = nodes; each cell = subproblem
 - Longest Increasing Subsequence $O(n^2)$:
 - $S[i] = LIS[i..n]$ starting at A[i]
 - $S[i] = (\max \text{ connected node after}) + 1$; $O(i) = O(n)$; for n nodes
 - Lazy Prize Collecting (max points k steps) $O(kE)$:
 - Transform G \rightarrow DAG w/ k copies per node, longest path per source
 - $P[v, k] = \max(P[w, k-1] + e(v, w) \text{ per } v.\text{nbrs})$
 - Min Vertex Cover on Tree (min set of nodes touching every edge) $O(V)$:
 - $S[v, 0/1]$ = size of VC in v subtree if v not/covered (2V subproblems)
 - $S[v, 0]$ = sum of S[child, 1]; $S[v, 1] = 1 + \text{sum of } S[\text{child}, 0]$ $O(1)$
 - All-Pairs Shortest Path with q queries $O()$:
 - Dijkstra per query: $O(qE \log V)$; w/ memoisation: $O(VE \log V)$
 - Preprocessing: APSP + memoise all, then $O(q)$ for queries
 - Weightless/same weights: $BFS * V = O(VE)$
 - Floyd-Warshall: $S[v, w, P8] = \min(S[v, w, P7], S[v, 8, P7] + S[8, w, P7])$
 - Consider only n+1 sets P in increasing size $O(V)$
 - Calculate SP from v-w: $O(V^2) \rightarrow O(V^3)$
 - Path reconstruction: store first hop/intermediate node
 - Transitive closure: return path adjacency matrix
 - Min bottleneck edge: return heaviest edge on path matrix

- MaxHeap**
 - Unsorted binary tree: keys = priorities (p) \rightarrow can be used to implement PQ
 - Invariant: $u.p \geq \max(u.\text{left}.p, u.\text{right}.p)$
 - Complete tree: every level full except leaves, leaves as far left as possible
 - bubbleUp: swap w/ parent; bubbleDown: swap w/ max child $O(\log n)$
 - Insert $O(\log n)$: ins @ end + bubbleU
 - extractMax/delete $O(\log n)/O(h)$: swap(last, deleted) + bubbleD
 - Represent: tree or compact array (seq: BFS order; variable-length)
 - $u.L = 2i$, $u.R = 2i+1$, $u.\text{parent} = \lfloor i/2 \rfloor$ OR binary: append 0/1 & remove last
 - Heapify: reverse traverse array & bubbleD if needed $O(n)$
 - Invar: all subtrees starting at [i..n] are heaps
 - HeapSort: heapify $O(n)$ + extract n times $O(n \log n) = O(n \log n)$